

# Programming Assignment 1: GNN-based Node Classification on Cora

Summary of Findings for Graph Learning (2025W)

*Simon Schindler, Marlene Grabner, Aryan Rahbari*

December 15, 2025

## 1 Introduction

This report summarizes the findings for the semi-supervised node classification tasks on the Cora citation network, as required by Programming Assignment 1. We utilized the PyTorch Geometric library for all experiments.

### Python Environment Setup

To run the Jupyter notebooks and reproduce the results, a specific Python environment is required. The following steps outline the setup process using uv, a fast Python package installer.

1. **Install uv:** If not already installed, run the following command:

```
curl -Lsf https://astral.sh/uv/install.sh | sh
```

2. **Create Virtual Environment:** This project uses Python 3.13. From the project root, create a virtual environment:

```
uv venv
```

uv will automatically detect the required Python version from the `.python-version` file.

3. **Install Dependencies:** Activate the environment and install packages using the `sync` command:

```
uv pip sync
```

This command uses the `pyproject.toml` and `uv.lock` files to ensure a reproducible environment. You can then start Jupyter Lab by running `jupyter lab`.

## 2 Task 2(a): Hyperparameter Optimization and Best Model

The objective of this task was to find a GCN-based model that achieves at least 80% test accuracy on the Cora dataset using the standard validation split. A comprehensive grid search was performed over key hyperparameters, including hidden dimensions, number of layers, dropout rate, learning rate (LR), and convolutional layer type (GCNConv, GraphConv, GATConv).

Table 1 summarizes the hyperparameters evaluated.

Hyperparameter	Evaluated Values
Hidden Dimensions	<b>16</b> , 64, 128
Number of Layers	<b>2</b> , 4, 8
Dropout	0.3, <b>0.5</b> , 0.7
Learning Rate	0.01, 0.005, <b>0.001</b>
Epochs	10, 50, <b>200</b>
Convolution Layer	<b>GCNConv</b> , GraphConv, GATConv

Table 1: Hyperparameters evaluated in the grid search. The best configuration is marked in **bold**.

## 2.1 Optimal GNN Architecture

The best-performing model, selected based on the highest validation accuracy, significantly exceeded the target, achieving a Test Accuracy of 88.50%. The optimal hyperparameters and resulting performance metrics are summarized in Table 2.

Table 2: Optimal GCN Architecture and Performance on Cora

Hyperparameter	Setting	Notes
Convolutional Layer	GCNConv	Standard GCN
Number of Layers	2	Shallow architecture
Hidden Channels	16	Smallest dimension tested
Dropout Rate	0.5	Optimal regularization
Learning Rate (LR)	0.001	Slow, stable convergence
Epochs	200	Full training run
Performance	Accuracy	
Training Accuracy	95.36%	
Validation Accuracy	88.00%	(Used for selection)
Test Accuracy	88.50%	(> <b>80%</b> target reached)

## 2.2 Key Hyperparameter Influences

Analysis of the grid search results revealed clear patterns regarding model performance and overfitting:

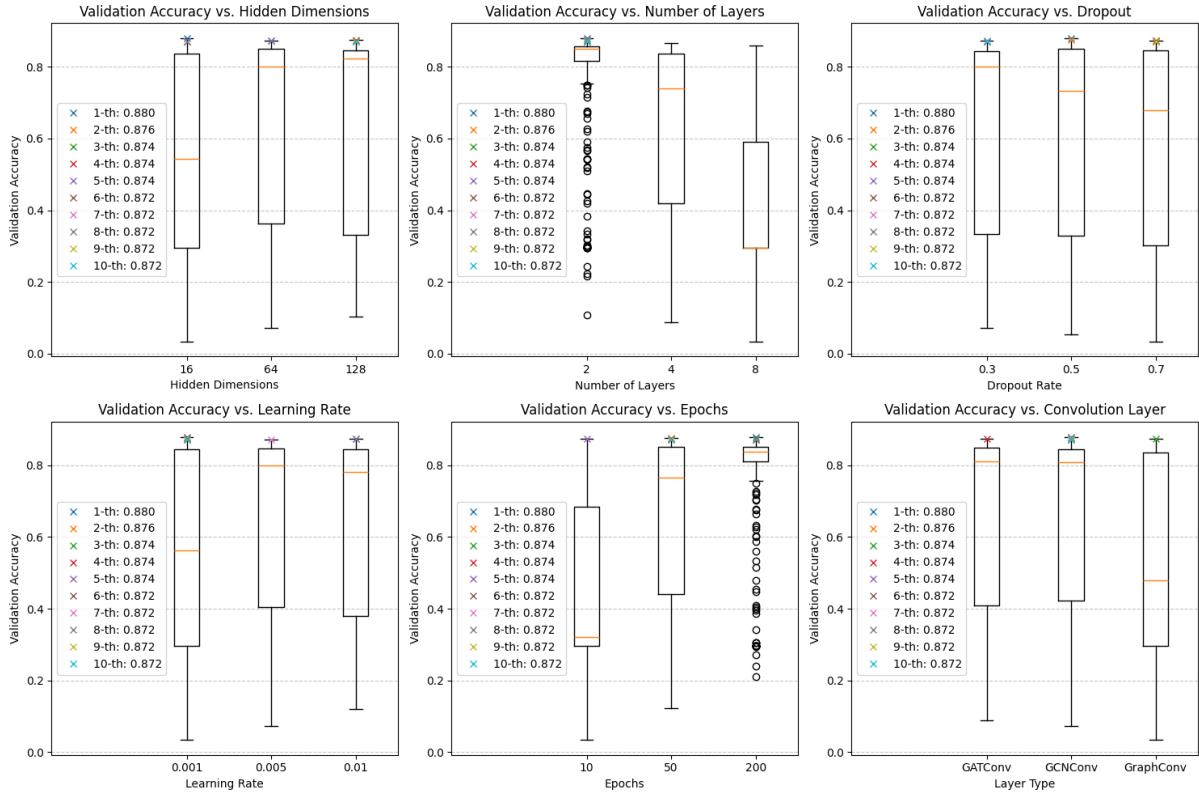


Figure 1: Influence of model hyperparameters on validation accuracy. Best 10 runs are individually marked with crosses.

Analyzing fig. 1 yields the following insights:

- **Number of Layers:** Performance peaked dramatically at 2 layers. Deeper models (4 and 8 layers) showed a significant drop in accuracy, indicating that the standard GCN architecture quickly suffers from **oversmoothing** on this dataset, where node representations become indistinguishable across the graph.
- **Hidden Dimensions:** The smallest dimension tested, 16, was optimal. Models with 64 or 128 channels showed decreased accuracy and a greater tendency toward overfitting (larger gap between training and validation accuracy).
- **Learning Rate:** A low LR of 0.001 was crucial for achieving the highest accuracy, suggesting that fine-tuning weights over many steps (200 epochs) is more effective than rapid learning.
- **Layer Type:** The standard GCNConv performed marginally better than GraphConv and significantly better than GATConv, suggesting the attention mechanism was not necessary for performance improvement in this setting.

### 3 Task 2(b): Homophily Analysis

This task focused on analyzing the graph's homophily and investigating the hypothesis that nodes with low homophily scores are more likely to be misclassified by the optimized GNN.

#### 3.1 Homophily Measurement

##### 3.1.1 Overall Homophily

The overall Node Homophily  $H_{\text{node}}$  was calculated using the formula:

$$H_{\text{node}} = \frac{1}{|\mathcal{V}|} \sum_{v \in \mathcal{V}} \frac{|\{(w, v) : w \in \mathcal{N}(v) \wedge y_v = y_w\}|}{|\mathcal{N}(v)|}$$

The computed homophily score for the Cora dataset is  $H_{\text{node}} \approx 0.8252$ . This confirms that Cora is a **highly homophilous graph**, where approximately 82.5% of a node's neighbors share its class label.

##### 3.1.2 Homophily Distribution

Figure /reffig:homohist visualizes a histogram of per-node homophily scores showing a severely skewed distribution, with the vast majority of nodes (over 65%) residing in the highest homophily bin (0.9 – 1.0), further solidifying the homophilous nature of the graph.

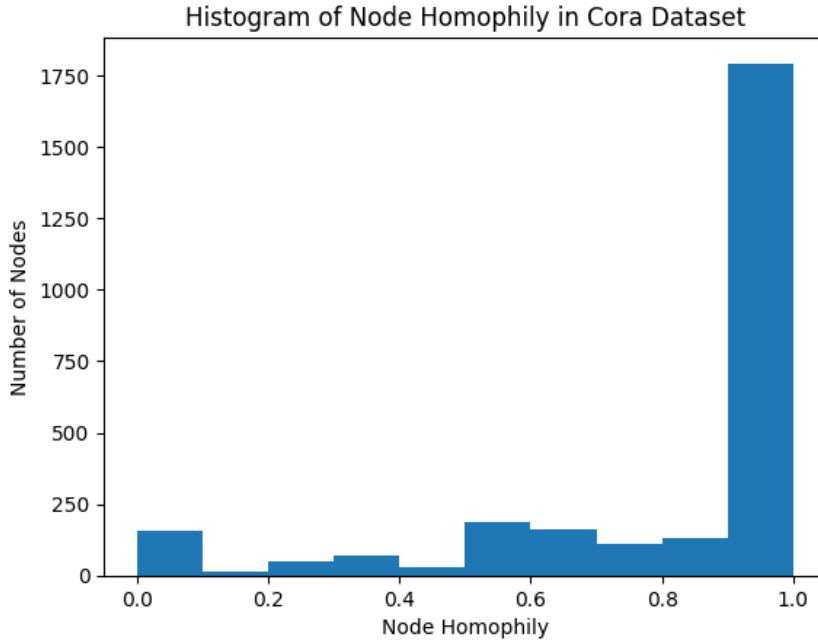


Figure 2: Homophily Distribution in Cora

#### 3.2 Predictive Performance vs. Homophily

##### 3.2.1 Experimental Design

To test the relationship between homophily and predictive performance, the nodes in the validation set were partitioned into 10 bins based on their individual homophily score (e.g., 0.0 – 0.1, 0.1 – 0.2, ..., 0.9 – 1.0). The classification accuracy of the top 10 models (identified in Task 2a) was then measured for the nodes within each bin.

### 3.2.2 Findings

The results, visualized in a box plot (Figure 3), provide a clear answer to the central question:

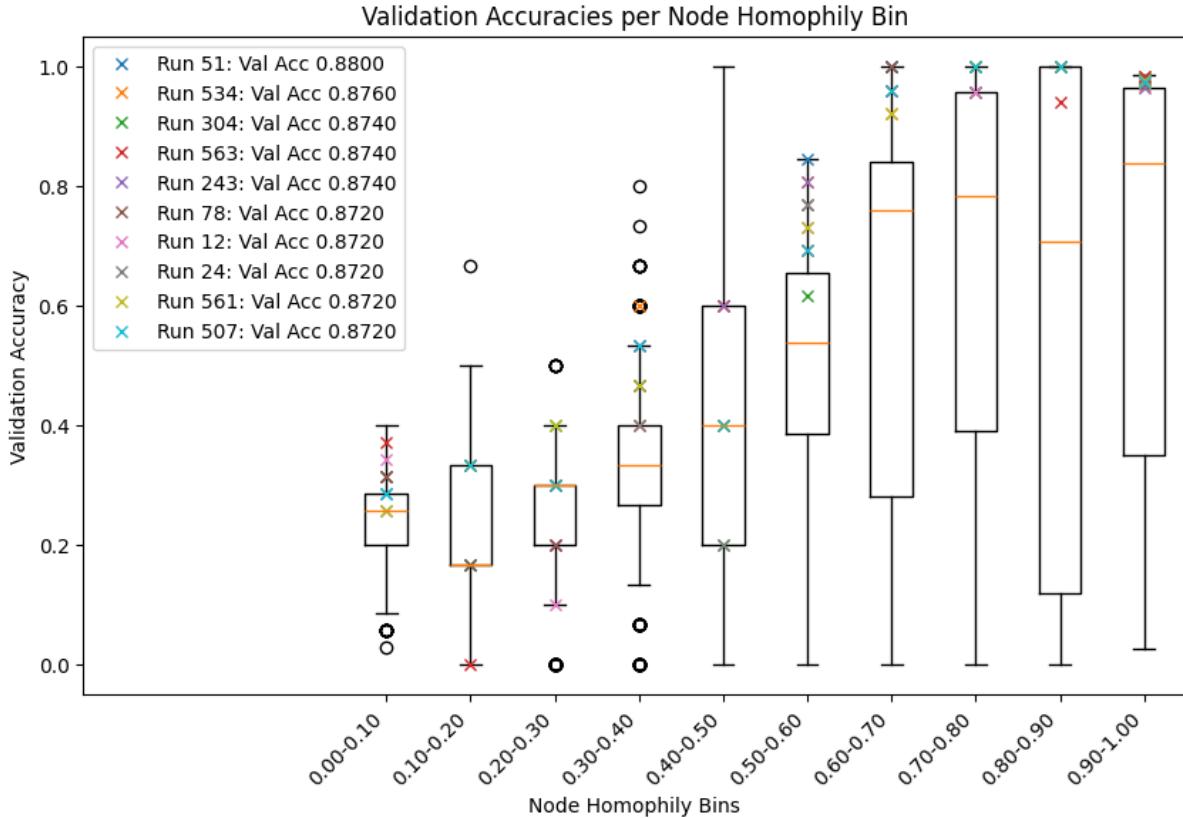


Figure 3: Validation Accuracy Distribution Across Node Homophily Bins

- **High Homophily ( $\geq 0.8$ ):** Models show **excellent performance**. In the  $0.9 - 1.0$  bin, the median accuracy across all runs approached 100%. The GCN performs optimally when the neighborhood signal is consistent.
- **Low Homophily ( $< 0.5$ ):** The models consistently struggle. The median accuracy in the  $0.0 - 0.1$  bin is very low (well below 20%), indicating that the majority of nodes in this bin are misclassified.
- **Conclusion:** Nodes with low homophily are **significantly more likely to be misclassified**. This is a characteristic failure mode of standard GCNs, which operate under the assumption of homophily. When a node is heterophilous (linked to different classes), the message-passing mechanism aggregates conflicting labels, corrupting the node's representation and leading to incorrect predictions.